

Parallel computing for 2D/3D meshing manipulation



Paul LAFOIX-TRANCHANT, Antoine OLEKSIK

Lyon university

Supervised by Fabrice JAILLET and Florence ZARA



Context

- 2D / 3D Mesh refinement (Region Of Interest)
- Real time simulation

Objective

- Analyze the existing library
- Find data structure to optimize parallelization
- Speed up the meshing process


Libraries for C++ parallelism

- IntelTBB
- + Thread-safe containers
- Only few examples

Mesh construction

1. Compute the Bounding Box

(i.e. create the minimal number of initial quadrants that contains the input surface)




2. Generate Quadrants


Process :

Repeat until desired Refinement Level
For each Quadrant :
If the quadrant intersects the chosen region
Split the quadrant into 4 new identical quadrants.
Remove those which are outside.


Splitting process :
in red and blue new edges
in green new points



Split all region




Split surface region



3. Create a balanced structure


(i.e. at most 1 subdivision level between two adjacent quadrants)
For each Quadrant :
if the quadrant has a subdivided edge (midpoint != 0)
check if the subdivided edge are subdivided

Balanced with surface region





4. Apply the transition patterns

(i.e. to produce a conformal mesh)




The transition patterns

Transition patterns applied on output of the meshing on surface region



OpenMP

+ Pre-processing directives



Mutex version

- Protect the critical regions with different mutex.
- Use of concurrent data structures, such as `tbb::concurrent_unordered_set` to replace the set of QuadEdges.

Critical regions

- read/write in concurrency in the set of quadedge = one edge could be split twice
- Write new points and new quadrants

Problems


- It is not possible to modify an element in a set (eg. In the set of QuadEdges) because of the hashtable used to sort elements. However, it is possible to modify a mutable attribute of an element if it's not used for sorting.

Results

Refinement applied on ALL Quadrants

At the end :


4,095,572 points, 13,653,086 QuadEdges and 4,057,578 Quadrants



Refinement applied on SURFACE Quadrants

At the end :

106,819 points, 320,520 QuadEdges and 55,533 Quadrants

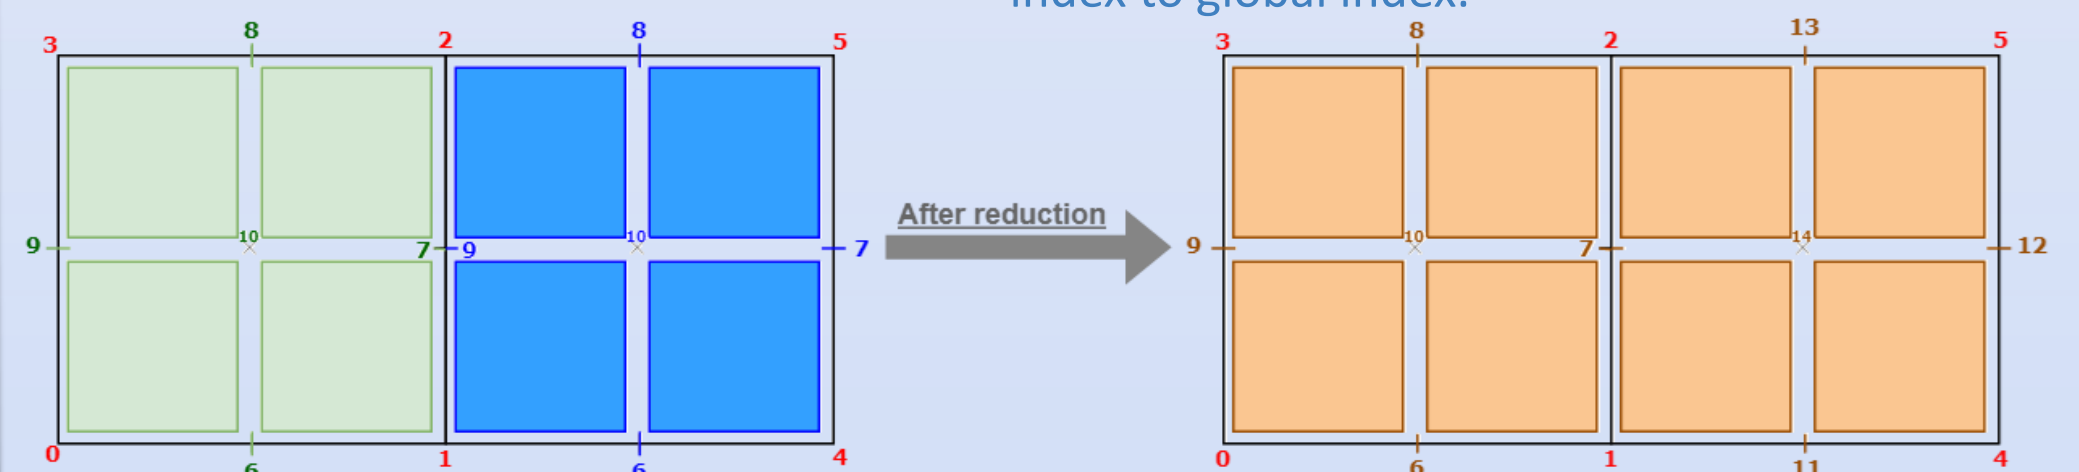


Parallelize the mesh construction



```
for (int rl=0; rl<desiredLevel; ++rl) {  
    #pragma omp parallel  
    {  
        int tn = omp_get_thread_num();  
        accumulation(newPts[tn], newEdges[tn], newQuad[tn]);  
    }  
    //Single thread  
    reduction(newPts, newEdges, newQuad);  
}
```

Reduction version

- No more critical regions.
- All threads have their own copy of Edges, Points and Quadrants filled in the accumulation part.
- The reduction part is done by a single thread and creates the final Edges, Points and Quadrants.
- Detection of identical points and update the local index to global index.



Before/after the reduction process. In red points before the accumulation done by green and blue threads. Notice that they start creating points at index 6 (the number of points before the accumulation)



Conclusion

- The mutex version is more performant than the reduction, even for high number of quadrants ($> 10^9$). The reduction part is done sequentially and should need optimizations
- Implementations avoiding critical sections produce a lot of new code, and it makes the whole project less maintainable
- Parallelism is not easy !

Perspectives

- Build a quadtree structure to represent all the Quadrants
- Speed up the reduction part of the version with reduction with a better algorithm and/or the use of parallelism
- Make the reduction process only when the desired refinement level is reached

